

Coherent Errors: Chess Transformers Act on Misremembered Boards

Mohit Pratap Singh Rathore
Oviqo Private Limited

Gunveer Sinh Kalsi
Oviqo Private Limited

Sriharsha Meduri
Oviqo Private Limited

Abstract

Long-horizon failure is reported as one number, which confounds two failures with opposite remedies: the model has lost track of the world, or it holds the world and chooses badly. We separate them where an exact state oracle exists, chess from move notation with no search at inference, in models of 3,440,064 to 39,884,288 parameters, reading linear probes against the true position at every token.

Aggregate state fidelity, the quantity usually reported, is the weaker predictor of an illegal next move within a fixed depth, at AUC 0.538 to 0.671 against 0.592 to 0.687 for the same representation restricted to the squares the move touches. And the errors are coherent. Reconstructing the position the model appears to believe, 0.355 of its illegal moves are legal there, against 0.050 for the same move under another position’s belief and 0.007 for an arbitrary illegal move. Editing the probe direction encoding an occupied square removes 0.934 of the probability mass on moves from it, a norm-matched random edit 0.077, so the representation is read and not merely present. Belief consistency rises monotonically across six model sizes; the causal edit is not measured at every size. A substantial identified subset of long-horizon errors is coherent with a degraded state representation that the policy demonstrably reads.

We report failures too: a backwards layer profile, a degenerate pre-registered metric, an auxiliary objective improving the measurement of state more than its consequences, and a second domain needing seven designs to become measurable.

Keywords: state tracking, world models, probing, model diagnostics, chess, long-horizon reasoning

1 Introduction

A model that reasons over many steps has to keep track of something. An attention-only decoder has no explicit bounded mutable state in which to keep it. Each token’s representation is rebuilt from the whole prefix through attention (Vaswani et al., 2017), so any quantity the task requires, a board position, the value of a variable, which sub-goal remains open, exists only because the network re-derives it, at every layer and every position, from the beginning. That re-derivation is dependable over short spans and unreliable over long ones. Anyone who has watched a model reason for forty steps knows the symptom: it stays fluent and quietly stops being right about its own earlier commitments.

The usual remedies sit outside the model. Tree search grafts deliberation onto the policy (Yao et al., 2023). Process reward models supervise intermediate steps with a learned verifier (Lightman et al., 2024). Self-consistency samples several chains and votes (Wang et al., 2023). Each raises end-task accuracy, and each leaves the original question untouched: what fails inside the network as the chain lengthens? The question matters commercially as well as scientifically, because the two candidate failures recommend opposite purchases. If the model has forgotten the world, a verifier

that rescores its outputs is paying repeatedly to repair something the architecture gave away. If the model remembers the world and reasons badly over it, a verifier is exactly right.

Answering the question requires reading an intermediate representation against a known-correct intermediate state, and on open-ended text no such state exists. The intermediate “state” of an argument in prose is not well defined, which is why evaluation in this area is noisy and why claims about it are hard to falsify. Our response is to move the measurement to a domain where the state is exact, cheap, and external, establish what can be learned there, and be explicit about what that does and does not license.

What we find. Working with search-free autoregressive models trained only on chess move notation, we measure the internal state against a rules engine at every token and compare it with what the model actually does. Three results carry the argument, and the rest of the paper tests how far they reach.

1. **The errors are coherent.** We reconstruct the board the model appears to believe and ask whether its illegal move is legal there. It is, for 0.355 of illegal moves, with a bootstrap interval of $[0.346, 0.362]$. The same move scored against another position’s believed board at equal depth is legal 0.050 of the time, and an arbitrary illegal move scored against the model’s own believed board is legal 0.007 of the time. Lossy decoding caps what is detectable at 0.654, the rate for moves that were genuinely legal.
2. **The representation is read, not merely present.** Editing the probe direction that encodes an occupied square removes 0.934 of the probability mass on moves from that square, while a norm-matched random edit removes 0.077 and shifts the mass at chance. The edit strength is calibrated rather than chosen.
3. **Where the world model is measured matters, increasingly with depth.** Whole-board probe accuracy predicts an illegal next move with an AUC of 0.538 to 0.671. The same representation restricted to the squares the move uses reaches 0.592 to 0.687. In the opening the two are almost indistinguishable; the gap opens as reasoning deepens, and in a joint fit the move-touched term carries a standardised weight of 0.504 against 0.118.

The first two are the ones we would defend hardest, and they belong together: the errors are coherent with a decayed state, and that state is causally load-bearing rather than an epiphenomenon. Failure here is not noise breaking through as the context grows. A substantial share of it is action admissible under a world that has drifted, which is a claim about state rather than about reasoning for that share, and it suggests why fluency outlives correctness: the model acts coherently with respect to a position it no longer holds.

Why chess. Chess presented purely as move notation, with no tree search at inference, gives what prose does not. The true state after any prefix is computable by a rules engine, so probing has a per-token label. Legality is decidable, so a state-tracking failure has an unambiguous observable rather than a judgement call. Move quality is separately scoreable by an engine, so the second failure mode can be measured conditional on the first being absent. Depth is uniformly scaled, unlike “number of reasoning steps” in prose. And the domain is combinatorial, so surface heuristics do not substitute for computation. Chess is our instrument, not our subject, and we are explicit in Section 9 about the reach that buys.

Contributions.

1. **A substantial identified subset of long-horizon errors is coherent.** Invalid actions are valid in the position the model appears to believe far more often than matched controls

allow, and the identified share rises across the range of model sizes we train.

2. **The state representation is read, not merely present.** Inducing a false belief by editing a single probe direction changes the action accordingly, at a calibrated edit strength and against a norm-matched control.
3. **Where a world model is measured matters.** Whole-board probe accuracy predicts whether the next action will be valid, and the same measurement restricted to the state that action consumes predicts it better, with the gap widening as reasoning deepens. We define a horizon on that fidelity directly, which avoids a failure mode of behavioural horizons that we document.
4. **A per-token diagnostic stratification** into state mismatch, belief-consistent error, conditional planning loss, and unresolved failure, measured against an exact oracle rather than a learned verifier, with the conditioning analysis that shows how far the strata can be told apart.
5. **An account of what failed**, including a degenerate metric, a prediction that ran backwards, an intervention that improved the measurement more than the behaviour, and seven task designs before a second domain became measurable. Each changed a decision and none is usually reported.

2 Related Work

Inference-time search. Deliberation over candidates (Yao et al., 2023), process supervision (Lightman et al., 2024), and sampling with a vote (Wang et al., 2023), built on chain-of-thought prompting (Wei et al., 2022), all buy accuracy with inference compute and take the base policy as given. Our concern is the base policy. A verifier that spends its budget repairing state-tracking errors is compensating for a representational defect at a recurring cost. What we do differently is diagnose which defect is present before recommending a remedy.

Long-context evaluation. Position-dependent retrieval failure (Liu et al., 2024) and effective-context measurement (Hsieh et al., 2024) show usable context is shorter than advertised context. Both measure retrieval of material present in the prefix. Long-horizon reasoning failure is different: the needed quantity was never written down, only implied, and has to be maintained through composition. A model can attend to every token of a sixty-ply game and still not know where the pieces are. What we do differently is measure maintenance rather than retrieval.

World models and probing. Linear probes recovering board state from game transcripts (Li et al., 2023), the linear structure of those representations (Nanda et al., 2023), chess as a state-tracking testbed (Toshniwal et al., 2022), and causal editing of internal associations (Meng et al., 2022) together establish that sequence models form internal world models and that intervening on them changes behaviour. Closest to our setting, Karvonen (2024) probes board state and estimated player skill in a chess-playing language model and intervenes on those representations, and the same board-state probes are later used as ground truth for evaluating dictionary learning (Karvonen et al., 2024). That work establishes the representation exists and is manipulable in this exact domain, which is our premise rather than our result. That literature answers whether a world model exists. We take its existence as given and ask a different question: how the fidelity of that world model varies with reasoning depth, and whether its failures explain the model’s behavioural failures at the level of individual actions. Probe design follows the control-task discipline of Hewitt and Liang (2019).

Concurrent work, and what is actually new here. The measurement at the centre of this paper is not new. Balogh and Jelasity (2026) report the fraction of illegal chess moves that are legal on a probe-reconstructed board, and reach a conclusion close to the opposite of ours: for models trained on large datasets that fraction is small, and the gradient of their board-state probe is nearly orthogonal to that of the next-token head, which they read as the probe lacking a causal role in generation. Li and Jiang (2026) separate state tracking from decision quality in searchless chess transformers, treating illegal moves as direct evidence of tracking failure. Walker and Lyons (2026) benchmark exact state tracking across architectures over a parameter range close to ours, and Pereira and Zuidema (2026) show world-state representations decaying during generation in frontier reasoning models, with a causal restoration experiment. Harang et al. (2025) document loss of coherent internal state over long sequences. State decay with depth, chess state tracking, and scaling this measurement are therefore not novel contributions of ours, and we do not claim them.

What differs is the conditions and the conclusion. Their illegal moves are adversarially induced; ours arise on ordinary held-out trajectories, which is a different population and may have a different mechanism. We add same-depth mismatched-belief and random-move controls, which bound how much of the effect a merely plausible board would produce. We compare aggregate against action-conditioned predictors on equal footing rather than assuming the aggregate is the right one. And where their gradient-orthogonality analysis finds the probe weakly coupled to generation, our calibrated residual-stream intervention finds a strong effect on the action. Those two results are not strictly contradictory, since a direction can be causally load-bearing at one layer while its gradient is near-orthogonal to the output head, but the tension is real and we flag it as the most useful open question this paper raises.

Expressivity of state. Theory bounds what these architectures can track: state-space models are weaker than their recurrent form suggests (Merrill et al., 2024), and intermediate tokens change what a Transformer can express (Merrill and Sabharwal, 2024). That work asks what is representable in principle. We ask what a trained model actually maintains, and where it stops.

Recurrence and state-space models. Structured state-space models (Gu et al., 2022) and selective variants (Gu and Dao, 2024) give constant per-token state, usually motivated by throughput and context length. We use a small supervised recurrent channel not as a contribution in itself but as a test case for whether the diagnostic discriminates between architectures.

Search-free chess models. Policies that reach strong play by distilling engine evaluations without a tree (Ruoss et al., 2024) show the task is learnable without search. Their objective is strength; ours is measurement, so we deliberately study models weak enough to fail. A model that never loses the position teaches nothing about losing it.

3 Design Problem and Requirements

Stated plainly: long-horizon failure is reported as a single number, so it cannot be acted on. A practitioner watching accuracy fall at step thirty cannot tell whether to buy a larger model, add a verifier, or change the architecture, because the measurement does not distinguish the causes.

The procedure rests on four commitments, each ruling out a way of fooling oneself. The diagnostic is grounded in an external oracle rather than model self-report. The two failure modes

are separated before any effect is attributed. Every decodability number is reported against trivial predictors that could produce it without a world model. And the representation is tested by intervention, not only by correlation.

In scope. Search-free autoregressive decoders; domains with an exact, cheaply computable state oracle; state of bounded description length.

Out of scope. Inference-time search, which we neither use nor evaluate; tasks whose state grows without bound; and claims about production-scale models, which our compute does not license.

4 Problem Formalisation

Let a task be a token sequence $x_{1:L}$ generated left to right, and $s_t = \Phi(x_{1:t})$ the exact state after the prefix, from an external oracle Φ . For chess $s_t = (b_t, \tau_t, c_t, e_t)$: piece placement, side to move, castling rights, and the en-passant target.

Our probe decodes only the first component. Write $\hat{b}_t^{(\ell)} = g^{(\ell)}(h_t^{(\ell)})$ for decoded piece placement, and let the pseudo-state used by the belief test be $\tilde{s}_t = (\hat{b}_t, \tau_t, \emptyset, \emptyset)$, with side to move taken from ply parity and castling and en-passant disabled so that nothing is licensed by an assumption we did not decode. Claims below concern $\hat{b}_t$ and $\tilde{s}_t$, not s_t . The state space is large, but the description length of any s_t is constant in t : a position occupies a few hundred bits however many plies produced it.

Let $h_t^{(\ell)}$ be the residual-stream activation at position t , layer ℓ , and $g^{(\ell)}$ a linear decoder trained to recover s_t from it. Write $\hat{s}_t^{(\ell)} = g^{(\ell)}(h_t^{(\ell)})$ for the decoded, or believed, state.

Definition 1 (Board-placement fidelity). $F^{(\ell)}(t) = \Pr[\hat{b}_t^{(\ell)} = b_t]$, the probability that piece placement is linearly decodable at depth t . We name it for placement rather than for state because side to move, castling rights and the en-passant target are not decoded.

Definition 2 (Move-touched fidelity). The probability that $\hat{s}_t$ and s_t agree on the origin and destination squares of the selected move. We name it for what it measures rather than calling it action-relevant, because it is a deliberately local proxy. The state an action actually depends on is larger: legality can turn on intermediate path squares for sliding pieces, the king’s square and post-move attack status, pins and discovered attacks, castling rights and transit squares, the en-passant target, side to move, and promotion. A fuller treatment would define $R_{\text{legal}}(a, s)$, the complete set of components legality depends on, and report a progression from whole-board through move-touched to path-aware and fully legality-aware fidelity. We report the local proxy and the controls that bound it, and leave that progression to future work.

Definition 3 (Behavioural error and its stratification). Let $Y_t = 1$ when the action selected at depth t is behaviourally invalid, and let $M_t = 1$ when the decoded state disagrees with the oracle on the components the selected action reads. By the law of total probability,

$$\Pr(Y_t=1) = \Pr(M_t=1) \Pr(Y_t=1 \mid M_t=1) + \Pr(M_t=0) \Pr(Y_t=1 \mid M_t=0).$$

This is an observational stratification and not a causal decomposition. The first term is not by itself memory-caused error and the second is not by itself policy-caused error.

Definition 4 (Belief-consistent error). Write $\mathcal{A}(s)$ for the actions valid in state s and $\tilde{\mathcal{A}}(\hat{s})$ for those admitted by the decoded state. An error is belief-consistent when $B_t = \mathbf{1}[a_t \notin \mathcal{A}(s_t) \wedge a_t \in \tilde{\mathcal{A}}(\hat{s}_t)]$, and we report $\Pr(B_t=1 \mid a_t \notin \mathcal{A}(s_t))$.




Figure 1: State decodability by layer and depth. Contrary to our prediction, the position is most linearly available at the final block rather than in the middle of the stack.

Two consequences of this framing are worth stating plainly, because an earlier version of this paper got them wrong. A state mismatch need not produce a behavioural error: our own measurements show the probe wrong on the squares a move touches for a substantial fraction of *legal* moves, so $M_t=1$ does not imply $Y_t=1$. And because the probe is an incomplete and lossy read of the state, belief consistency identifies a subset of errors rather than partitioning all of them. The errors it does not identify are *unresolved*: they may be genuine policy failures, state errors the probe did not capture, information the probe cannot read linearly, or reconstruction failures. We do not call them planning errors.

5 The State-Space Bottleneck

Proposition 1 (Cost mismatch). *Under full attention each new token reads an $O(t)$ key-value history and retains $O(t)$ cached memory, while the task state has description length $O(1)$. A key-value cache does amortise the computation of earlier activations, so the cost is not recomputation from scratch; what the architecture lacks is a bounded mutable state that could carry s_t forward and be updated in place. We offer this as a motivating cost mismatch, not as a proof that fidelity must decay.*

Two consequences were testable and we tested both. Fidelity should fall with depth even well inside the context window, because the re-derivation circuit must compose more updates within a fixed layer budget. It does (Section 7.1).

We also predicted fidelity would peak at intermediate layers and fall toward the output, the state being a means rather than an end. **That prediction is wrong.** Decodability rises through the stack and is maximal at the final block, from 0.323 at the first layer to 0.543 at the last. For this task the state sits close to the target: the next legal move is a direct function of the position, so a representation making the position linearly available is also a good representation for the output head. We keep the failed prediction because it constrains where a state carrier should be inserted, and our spacing was chosen before we knew.

Table 1: State decodability against every trivial explanation, for the baseline model. Occupied-square accuracy at the most decodable layer, beside a per-square per-depth majority predictor, a label-permutation control, and a randomised-weight model of identical architecture. The gap is what training put there.

Ply	model	majority	label perm.	random weights	illegal rate
0–10	0.975	0.867	0.879	0.898	0.0004
10–20	0.876	0.599	0.673	0.641	0.0115
20–30	0.761	0.444	0.487	0.439	0.0439
30–40	0.655	0.359	0.315	0.297	0.0982
40–50	0.557	0.261	0.166	0.195	0.1766
50–60	0.475	0.073	0.072	0.141	0.2347
60–80	0.353	0.000	0.019	0.071	0.2920
80–120	0.173	0.000	0.001	0.038	0.3231
120–161	0.066	0.000	0.000	0.030	0.3192

6 Diagnostic Protocol

Data. Public game records above a rating floor, deduplicated by move-sequence hash, split by *game*, so no position from a given game appears in more than one split. Splitting by game does not make the splits prefix-disjoint, since distinct games share openings, and we do not claim it does. The corpus holds 120,000 games: 96,000 for training, 12,000 for probes, 12,000 for evaluation, mean length 78.2 plies, move vocabulary 1,926.

Oracle. A rules engine replays every game and emits the exact state after every token, verified complete before any model is measured.

Probes. One linear head per layer, fitted on the probe split and reported on the evaluation split. Because most squares are empty, raw per-square accuracy is misleadingly high, so we report accuracy on occupied squares against a per-square, per-depth majority predictor.

Controls. A majority predictor; a label-permutation control pairing activations with another game’s state at equal depth, preserving marginals while destroying the relation to the input; and a randomised-weight model of identical architecture, bounding what probe capacity alone extracts.

Behavioural measures. Illegal top-1 move rate against depth, and the probability mass on legal continuations, both read against the oracle.

7 Results

7.1 State and behaviour both decay with depth

Occupied-square decodability falls from 0.975 in the opening to 0.066 in the deepest bucket, while the illegal top-1 rate rises from 0.0004 to 0.319 and the probability mass on legal moves falls from 0.993 to 0.455. At mid-game the trained model is well clear of every trivial explanation, reaching 0.555 against 0.261 for the majority predictor and lower still for the label-permutation and randomised-weight controls, which Table 1 reports in full at every depth. There is a real state

Table 2: Aggregate and move-touched state error compared like for like, as predictors of the same binary outcome on the same samples: whether the model’s top-1 move is illegal. AUC is computed within each depth bucket so neither predictor can simply be tracking depth. The standardised coefficients come from a logistic fit containing both predictors. The aggregate measure is informative and it is the weaker of the two, and the gap widens with depth.

Ply	AUC aggregate	AUC move-touched	β aggregate	β move-touched	n
10–20	0.671	0.687	+0.065	+0.153	15,000
20–30	0.627	0.681	+0.120	+0.332	15,000
30–40	0.584	0.662	+0.100	+0.506	14,511
40–50	0.548	0.643	+0.057	+0.573	13,176
50–60	0.551	0.617	+0.114	+0.500	11,262
60–80	0.538	0.592	+0.076	+0.431	16,426
pooled, depth as covariate	n/a	n/a	+0.118	+0.504	85,375

Table 3: Does state loss explain behaviour? The aggregate column is the within-depth AUC of whole-board probe error as a predictor of an illegal move, reported as a range over depth buckets. Move-touched error is the probability the probe is wrong on the origin or destination square. Belief consistency is the share of illegal moves admissible under the model’s own decoded board, beside the same move under another position’s belief, an arbitrary illegal move, and the legal-move recovery reference.

Model	Params	aggregate	move-touched		belief consistency			
		AUC	illegal	legal	own	mismatch	random	recovery
Attention-only	7,345,920	0.54 to 0.67	0.806	0.384	0.355	0.050	0.007	0.654
Attention-only, parameter-matched	7,505,976	n/a	0.804	0.391	0.350	0.051	0.007	0.650
Recurrent channel ($\lambda=0$)	7,508,416	n/a	0.801	0.384	0.358	0.053	0.007	0.658
Recurrent channel ($\lambda=1$)	7,508,416	n/a	0.806	0.399	0.366	0.093	0.021	0.656
6L/192	3,440,064	0.52 to 0.70	0.831	0.447	0.266	0.049	0.006	0.581
12L/384	22,835,328	0.53 to 0.71	0.747	0.327	0.412	0.056	0.009	0.713
12L/512	39,884,288	0.53 to 0.74	0.711	0.297	0.439	0.059	0.009	0.745

representation, and it becomes progressively less recoverable by a linear decoder. Section efsec:drift tests whether that reflects loss or a moving basis.

7.2 Aggregate fidelity is the weaker signal

The two curves fall together, and it would be easy to stop there and call them one phenomenon. The relationship is real but weaker than that, and measuring it carelessly overstates it.

An earlier version of this analysis compared a point-biserial correlation for the aggregate predictor against a conditional probability for the move-touched one. Those are different statistics, and the contrast flattered the conclusion: a correlation against a heavily skewed binary outcome is attenuated, so the aggregate measure looked uninformative when it was merely differently scaled. We now score both the same way, as predictors of the same binary outcome on the same samples, within depth so that neither can simply be tracking depth.

Aggregate board fidelity predicts an illegal next move with an AUC between 0.538 and 0.671. That is above chance and not negligible. The same representation restricted to the two squares the move touches reaches 0.592 to 0.687, and the two are closest in the opening and furthest apart

Table 4: Causal intervention at the calibrated edit strength. Editing the state direction removes probability mass from moves out of the affected square; a norm-matched random direction does not. The manipulation check reports how often the edit actually changed the decoded belief.

Model	n	belief flipped	mass before	mass after	random edit after
Attention-only	700	1.000	0.071	0.005	0.065
Attention-only, parameter-matched	600	1.000	0.074	0.004	0.066
Recurrent channel ($\lambda=0$)	600	1.000	0.070	0.004	0.064
Recurrent channel ($\lambda=1$)	600	1.000	0.067	0.004	0.062

Table 5: Horizon and state fidelity by condition. $H_{\text{ill}}(.05)$ is the ply at which the illegal-move rate first crosses five percent, interpolated between bucket midpoints because the bucketed read-off returns the same bucket for every condition and so discriminates nothing. Mean $\pm$ s.d. over seeds where more than one was run.

Model	Params	seeds	occ. acc. @40–50	illegal @40–50	$H_{\text{ill}}(.05)$
Attention-only	7,345,920	3	0.555 ± 0.003	0.176 ± 0.002	26.1 ± 0.0
Attention-only, parameter-matched	7,505,976	3	0.554 ± 0.002	0.172 ± 0.003	26.0 ± 0.1
Recurrent channel ($\lambda=0$)	7,508,416	3	0.556 ± 0.003	0.169 ± 0.002	26.2 ± 0.3
Recurrent channel ($\lambda=1$)	7,508,416	3	0.632 ± 0.003	0.166 ± 0.002	26.5 ± 0.2
read-out ablated at test	n/a	3	0.559 ± 0.005	0.149 ± 0.005	27.2 ± 0.4
6L/192	3,440,064	3	0.468 ± 0.006	0.245 ± 0.003	21.4 ± 0.1
8L/256	7,345,920	3	0.555 ± 0.003	0.176 ± 0.002	26.1 ± 0.0
12L/256	10,504,960	3	0.577 ± 0.005	0.149 ± 0.001	27.8 ± 0.2
8L/384	15,737,472	3	0.632 ± 0.006	0.134 ± 0.004	29.5 ± 0.2
12L/384	22,835,328	3	0.656 ± 0.001	0.112 ± 0.002	31.7 ± 0.3
12L/512	39,884,288	3	0.694 ± 0.003	0.101 ± 0.001	33.8 ± 0.1

at depth: at ply 40 to 50 the aggregate reaches 0.548 against 0.643. Fitted jointly with depth as a covariate, the move-touched term takes a standardised weight of 0.504 against 0.118 for the aggregate.

The claim we can defend is therefore narrower than the one we first wrote, and it is still a claim about a standard practice. Whole-board probe accuracy is a usable but weak proxy for whether a world model is load-bearing, it degrades as a proxy exactly where long-horizon behaviour matters most, and a measurement restricted to the state an action consumes is several times more informative from the same probe and the same activations. Stated as a conditional rather than a ranking, the probe is wrong on the squares a move touches for 0.806 of illegal moves against 0.384 of legal ones. The relevant question is not how much of the world the model has lost, but whether it has lost the part it is about to use.

7.3 Lost, or merely moved? Depth-local probes

Every fidelity number so far comes from one decoder fitted across all depths, and a falling number has been read as the state becoming less available. A single global decoder cannot distinguish that from an alternative: the state might stay equally present while its linear basis moves with depth, so that one fixed decoder simply transfers badly. The reversible-cycle experiment holds the position fixed but still uses the global probe, so it inherits the same ambiguity.

Table 6: The scale ladder, three seeds per rung. H_{ill} is read off behaviour, the ply at which the illegal-move rate crosses five percent. H_{fid} is read off the representation, the depth at which the probe first fails on the squares the chosen move touches more than half the time, and so cannot be moved by hedging. Belief consistency is normalised by each model’s own legal-move recovery reference, because that reference rises with size. Mean $\pm$ s.d. over seeds.

Model	Params	seeds	$H_{\text{ill}}(.05)$	$H_{\text{fid}}(.5)$	belief / reference	mismatch
6L/192	3,440,064	3	21.4 ± 0.1	33.8 ± 0.2	0.476 ± 0.017	0.051 ± 0.002
8L/256	7,345,920	3	26.1 ± 0.0	40.8 ± 0.5	0.534 ± 0.008	0.052 ± 0.002
12L/256	10,504,960	3	27.8 ± 0.2	42.1 ± 0.3	0.552 ± 0.007	0.051 ± 0.002
8L/384	15,737,472	3	29.5 ± 0.2	50.0 ± 0.4	0.568 ± 0.008	0.053 ± 0.001
12L/384	22,835,328	3	31.7 ± 0.3	52.5 ± 0.5	0.577 ± 0.005	0.053 ± 0.003
12L/512	39,884,288	3	33.8 ± 0.1	59.1 ± 0.5	0.587 ± 0.011	0.056 ± 0.004

Table 7: Planning error against how strictly we condition on state. Each cell gives the median engine centipawn loss, the blunder rate, and the sample size, over legal moves whose touched squares the probe recovers and with at most the stated number of other squares wrong. Tightening the conditioning lowers the loss, so part of what looks like degraded judgement is state loss elsewhere on the board. Loss still grows with depth at fixed conditioning, so part of it is not. Cells with fewer than thirty scored positions are reported as n/a, and that scarcity is itself the finding: at depth a nearly intact board is rare.

Ply	≤ 0	≤ 2	≤ 4	≤ 8	≤ 16	any
20–30	n/a	24 / 0.18 (115)	24 / 0.21 (300)	31 / 0.22 (300)	44 / 0.31 (300)	58 / 0.34 (300)
30–40	n/a	n/a	38 / 0.25 (85)	62 / 0.36 (300)	75 / 0.41 (300)	91 / 0.47 (300)
40–50	n/a	n/a	n/a	67 / 0.41 (300)	120 / 0.53 (300)	111 / 0.52 (300)

We fit a separate linear probe inside each depth bucket and evaluate every probe on every bucket.

The diagonal answers the question. A probe fitted for the deepest bucket and tested there reaches 0.340, against 0.910 for the shallowest bucket tested on itself. Depth-local probes collapse in almost the same way the global probe does, 0.907 to 0.409, so the decline is not an artifact of a fixed decoder failing to keep up. If the representation merely rotated, a probe fitted at the relevant depth would recover what a global probe misses, and it does not.

Basis change is nonetheless present in the off-diagonal. A probe fitted on the shallowest positions reaches only 0.151 on the deepest ones, far below what a probe fitted there achieves, so the encoding does move with depth. It moves and it also thins, and the thinning is what the paper’s curves track.

One caveat cuts in our favour rather than against. Each depth-local probe sees roughly a sixth of the data the pooled probe sees, which is why the pooled probe edges the diagonal at depth. More data would help the local probes, so this is a conservative test of the drift account: it would have to overturn a gap that already runs the wrong way for it.

7.4 Which squares, not how many

Move-touched fidelity scores two squares where the aggregate averages sixty-four, and a statistic over two concentrated variables can look sharper than one over sixty-four diluted ones for reasons unrelated to relevance. The control that separates those explanations is a random pair.

Two randomly chosen squares predict an illegal move at an AUC of 0.539, well below the

Table 8: Domain 2, permutation tracking. Answer accuracy at the shallowest and deepest depths against a chance rate of one in four, the probe’s recovery of the queried object at depth, and belief consistency beside its mismatched control. The attention-only model solves the task well above chance while its state stays at chance under the probe, so it is not maintaining a linearly decodable assignment. The supervised channel is, and its errors are coherent with it.

Model	seeds	acc. shallow	acc. deep	probe deep	belief / control
Attention-only	3	1.000 ± 0.000	0.589 ± 0.010	0.263 ± 0.010	0.256 ± 0.005 / 0.249 ± 0.003
Recurrent channel ($\lambda=1$)	3	1.000 ± 0.000	0.989 ± 0.001	0.983 ± 0.001	0.701 ± 0.017 / 0.254 ± 0.009




Figure 2: Internal state and external behaviour decay together. (a) Occupied-square state accuracy at the most decodable layer, against a majority predictor, a label-permutation control, and a randomised-weight model of identical architecture. (b) The illegal top-1 move rate and the probability mass the model places on legal continuations.

whole board at 0.629. Two squares drawn at random from those actually *occupied* reach 0.631, statistically indistinguishable from the whole board. The squares a move touches reach 0.696, with a game-clustered interval that does not overlap either control. The advantage is therefore about which squares are read, not how many, and it is not an artifact of averaging fewer variables. That control equalises count and occupancy only in part: a move’s origin is occupied by the side to move, while its destination is usually empty or held by the opponent, so a pair of occupied squares does not reproduce the endpoint profile exactly. A fully matched control would draw one square occupied by the mover and one matching the destination’s true class.

Splitting the pair, the origin square carries more of the signal than the destination, 0.664 against 0.637. That ordering is what one would expect if the dominant failure is proposing a move from a square the model has misremembered as occupied.

7.5 The errors are coherent

If state loss drives the failure, an invalid action should be valid in the state the model believes it is in. We reconstruct that believed position from the probe and re-score the model’s own move in it, clearing castling rights so nothing is licensed by an assumption we did not decode, and testing pseudo-legality because a decoded board may lack a king.

Of the illegal moves the model plays, 0.355 are legal in its own believed position, with a game-




Figure 3: (a) Within a fixed depth, the correlation between the number of misremembered squares and playing an illegal move is small. (b) The same representation, scored only on the squares the move actually uses, separates illegal from legal moves sharply. Table 2 compares the two on equal footing, as predictors of the same outcome.

Table 9: Is the state lost, or is the representation rotating? Each row is a linear probe fitted inside one depth bucket and evaluated on every bucket; the pooled row is a probe fitted across all depths, as used elsewhere in the paper. Depth-local probes do not hold up at depth either, so the decline is not explained by a fixed decoder failing to transfer. Local probes see roughly a sixth of the pooled training data, which makes this a conservative test of the drift account.

Probe fitted at	ply 10	ply 20	ply 30	ply 40	ply 50	ply 60
ply 10	0.910	0.776	0.521	0.307	0.203	0.151
ply 20	0.835	0.789	0.695	0.540	0.348	0.201
ply 30	0.640	0.736	0.671	0.610	0.546	0.412
ply 40	0.253	0.591	0.636	0.553	0.486	0.416
ply 50	0.106	0.341	0.527	0.514	0.430	0.346
ply 60	0.097	0.241	0.415	0.497	0.447	0.340
pooled across depths	0.907	0.804	0.694	0.591	0.506	0.409

clustered interval of $[0.346, 0.364]$ over 13,802 illegal moves from 1,155 games. Positions within a game are not independent, so the interval resamples whole games. Two controls bound the alternatives. The same move scored against a believed position decoded elsewhere at equal depth is legal 0.050 $[0.046, 0.053]$ of the time, so the effect is specific to this belief at this position rather than to plausible boards generally. An arbitrary illegal move scored against the model’s own believed position is legal 0.007 $[0.006, 0.009]$ of the time, so the believed board is not permissive. For reference, genuinely legal moves are recovered as legal in the believed board 0.654 of the time, which is below one because probe decoding is lossy. We call this a legal-move recovery reference rather than a decoding reference, and we do not divide by it to estimate an unobserved true rate: that would assume probe sensitivity is the same on legal and illegal cases, which we have not established. What the comparison licenses is that a substantial identified subset of illegal moves is consistent with the decoded state, not that most of them are.

We do not divide by that reference. Doing so would estimate an unobserved true rate only under an assumption we have not tested, that the probe is equally sensitive on legal and illegal

Table 10: Which squares carry the signal. Each row scores probe error over a different set of squares as a predictor of an illegal top-1 move, on identical samples. Intervals are bootstrapped over games rather than positions, since positions within a game are not independent. Two random occupied squares are a partial match to a move’s endpoints, the origin being occupied by the mover while the destination is usually empty or held by the opponent, so this control equalises count and occupancy only in part. It performs like the whole board; the squares the move actually touches do not, so the advantage is about which squares rather than how many.

Squares scored	AUC [95% CI]
whole board, 64 squares	0.629 [0.623, 0.633]
two random squares	0.539 [0.534, 0.542]
two random occupied squares	0.631 [0.627, 0.636]
origin square only	0.664 [0.660, 0.668]
destination square only	0.637 [0.633, 0.642]
move-touched, origin and destination	0.696 [0.693, 0.700]

cases. What the comparison licenses is narrower and still substantial: a large identified subset of illegal moves is admissible under the board this model appears to hold at this position, at 7.2 times the rate of the matched control. It identifies state-consistent failure. It does not assign the remainder.

7.6 Inducing a false belief changes the action

Belief consistency is observational. It shows the model errs coherently with a decayed state without showing that the state representation drives the action. We therefore intervene. At a position where the side to move genuinely occupies a square, we push the residual stream along the probe direction encoding "this square is empty" and ask whether the model stops proposing moves from it.

Two safeguards decide whether such an experiment means anything, and both are reported. The manipulation check asks whether the edit changed the decoded belief at all: at the strength we use it does, in 1.000 of cases. The norm-matched control applies a random direction of identical magnitude at the same position, which separates belief from mere perturbation sensitivity.

Edit strength has to be calibrated rather than chosen, because a large enough push degrades the forward pass and will move behaviour for reasons having nothing to do with state. Table 11 reports the sweep. Below a quarter of the residual-stream magnitude the edit does not reliably flip the belief; above roughly one, the random control also starts destroying the output distribution, so any effect there is damage rather than belief. We report the intervention at 0.50, inside the window where the manipulation succeeds and the control does not.

At that strength, editing the state direction removes 0.934 of the probability mass on moves from the affected square, decreasing it in 0.996 of 700 cases. The norm-matched random edit removes 0.077 and decreases the mass in 0.484 of cases, which is indistinguishable from a coin flip. The state representation is not a decorative correlate of the model’s computation. It is read.

7.7 How much of the planning error is really memory

The other half of the decomposition needs care, because the conditioning that would isolate it cleanly is unattainable. Requiring the probe to recover the exact board leaves no samples past the opening, for the reason already given. Requiring only the squares the move touches leaves samples




Figure 4: The errors are coherent. (a) Against depth, restricted to buckets with at least thirty illegal cases. (b) Pooled: an illegal move is far more likely to be legal in the model’s own believed board than under another position’s belief at equal depth, or than an arbitrary illegal move is under the model’s own belief. The dashed line is the recovery reference for this test, set by how often genuinely legal moves are recovered as legal.

everywhere but does not constrain the rest of the board, so state loss elsewhere is still counted as a failure of judgement.

We therefore vary the strictness and read the trend (Table 7). Tightening the conditioning lowers the loss substantially. At ply 30 to 40 the median centipawn loss falls from 91 with the board unconstrained to 38 when at most four other squares may be wrong, and the blunder rate from 0.473 to 0.247. Roughly half of what a naive reading would call degraded judgement is misattributed memory failure.

A limit of this analysis is worth naming, because it is not a nuisance but the same fact the paper is about. The strictest cells are thinly populated at depth however many games we scan, since a nearly intact board becomes rare exactly where we would most like to condition on one. Cells with fewer than thirty scored positions are reported as n/a rather than given a number.

The trend does not go to zero. Holding the conditioning fixed at four wrong squares, the median loss still grows with depth, from 24 at ply 20 to 30 to 38 at ply 30 to 40. Judgement given intact state is not depth-invariant.

We take the honest reading. The identified state-consistent component is substantial, stricter state conditioning reduces measured planning loss, and a residual depth-dependent planning component remains. Calling memory dominant would overstate what a lossy probe and small conditioned samples support. Our original framing of H4, a clean separation in which an intervention moves one term and leaves the other untouched, was too tidy for what the measurement supports. The decomposition still earns its place, because it converts one uninterpretable curve into two quantities with different remedies, but a practitioner should expect the terms to be entangled rather than orthogonal, and should report the conditioning they used when they quote either one.

7.8 Scale

The behavioural horizon is read off the illegal-move curve at the ply where it first crosses five percent, interpolated between bucket midpoints because the bucketed read-off returns the same bucket for every condition and so cannot discriminate between them. A second horizon defined on




Figure 5: Choosing the edit strength. Too small and the belief does not flip; too large and a norm-matched random direction also destroys the output distribution, so any effect there is damage rather than belief. We report inside the shaded window.

the representation is introduced in Section 7.16; both are tabulated together there.

Across the ladder the behavioural horizon grows with scale and grows slowly, from 21.4 ply at 3,440,064 parameters to 33.8 at 39,884,288, three seeds per rung. Roughly an order of magnitude in parameters buys about half again as much depth, which is consistent with H2 and with the cost mismatch of Section 5: scaling helps, and it does not buy back a missing mechanism cheaply.

The question that matters more for the reach of this paper is not whether the horizon moves with scale, which it does, but whether the mechanism does. Figure 6(b) and Table 6 report belief consistency at every rung, three seeds each. The raw rate rises with size, but so does each model’s legal-move recovery reference, so the honest quantity is the ratio of the two. Normalised that way the coherence of errors still rises monotonically, from 0.476 at 3,440,064 parameters to 0.587 at 39,884,288, with seed spreads of a few parts in a thousand, while the mismatched control stays flat near 0.050 and the random-move control near 0.007. We show the ratio as a descriptive normalisation, since the reference itself rises with size, and not as an estimate of a hidden true proportion. What we have characterised strengthens with scale over the range we can reach rather than washing out, which is evidence against its being an artifact of the smallest models.

Both horizons order the models identically and both grow sublinearly in parameters, the behavioural one from 21.4 to 33.8 ply and the fidelity one from 33.8 to 59.1 (Table 6). Roughly an order of magnitude in parameters buys about half again as much depth, which is consistent with H2 and with the cost mismatch of Section 5: scaling helps, and it does not buy back a missing mechanism cheaply.

7.9 An intervention, evaluated by the diagnostic

The diagnostic is only useful if it discriminates. We test a fixed-width gated recurrent channel inserted after every k -th attention block, carrying $z_t \in \mathbb{R}^{d_s}$ with $d_s \ll d$:

$$u_t = W_{\text{in}} h_t, \quad \alpha_t = \sigma(W_g[h_t; z_{t-1}]), \quad z_t = (1 - \alpha_t) \odot A z_{t-1} + \alpha_t \odot \phi(u_t),$$

read back as $h_t \leftarrow h_t + W_{\text{out}} z_t$, with A diagonal so inference-time state is $O(1)$ per token, and an optional auxiliary loss $\lambda \mathcal{L}_{\text{state}}$ supervising z_t to be the state. We separate architecture ($\lambda =$




Figure 6: (a) The horizon against model size. (b) Belief consistency and its matched control against model size. The question is whether the mechanism is a property of small models or survives scale.

0), supervision ($\lambda > 0$), and read-out (zeroing W_{out} at test time), against a parameter-matched attention-only control whose feed-forward block absorbs the same parameter count.

The recurrent channel gives a clean and slightly uncomfortable answer.

With the auxiliary objective active, state decodability rises sharply, from 0.555 for the attention-only baseline to 0.632 at matched parameters and matched tokens, a difference many times the seed spread. The architecture without the objective, 0.556, is indistinguishable from the baseline, as is the parameter-matched control at 0.554. So the gain comes from supervising the channel, not from adding it, and not from the parameters it costs.

Behaviour barely follows. The illegal-move rate improves from 0.176 to 0.166, and the horizon from 26.1 to 26.5 ply. Set that beside the scale ladder: the attention-only model at 15,737,472 parameters reaches almost exactly the same state decodability, 0.632, and converts it into a much larger behavioural gain, 0.134 and 29.5 ply.

Two models with the same measured world model, then, and very different behaviour. This is the paper’s thesis arriving from the other direction. In Section 7.2 aggregate decodability predicted behaviour weakly across positions within a model; here it fails to predict behaviour across models. An auxiliary loss that makes the state linearly available does exactly what it is asked to do, and being available is not the same as being used.

An ablation that improves behaviour. Zeroing the read-out at test time lowers the illegal-move rate, from 0.166 to 0.149, and lengthens the horizon, from 26.5 to 27.2 ply. Removing a component the model was trained with should not help, so we measured what else it does. Language-model loss rises when the read-out is removed, from 3.010 to 3.199 across seeds, and the read-out is not vestigial: it injects a vector of 0.150 of the residual stream’s own norm. The ablated model is worse at predicting the continuation and simultaneously better at staying legal, and the reason is visible in the distribution. It places more mass on legal moves than the intact model, 0.647 against 0.615, while committing less confidently to any particular one.

The lesson is about our own observable rather than about the artifact. A model can lower its illegal-move rate by hedging toward the legal-move prior, so the illegal rate is not a monotone proxy for the quality of state maintenance, and should be read alongside a likelihood. This does not disturb the coupling, belief-consistency, or causal results, which are computed per position rather than from a rate, but it does qualify the intervention comparison and the horizon metric, and we

Table 11: Calibrating the edit strength for the causal intervention. Strength is in units of the residual stream’s own magnitude. The usable window is where the manipulation check succeeds while the norm-matched random control still behaves like a coin flip. Values from `results/attn_8L256_s0_patch_sweep.json`.

strength	belief flipped	state edit: mass after	state decreased	random: mass after	random decreased
0.10	0.652	0.046	0.980	0.076	0.540
0.25	0.996	0.020	0.988	0.078	0.452
0.50	1.000	0.005	0.996	0.065	0.484
1.00	1.000	0.000	0.996	0.054	0.556
2.00	1.000	0.000	0.996	0.025	0.712
6.00	1.000	0.000	1.000	0.013	0.784

flag it rather than leaving a reader to find it.

We therefore report H3 as supported only in its weakest form and H4 as not supported in the form we predicted. We had expected an intervention that improved state tracking to improve behaviour through it. What we observe is an intervention that improves the measurement of state tracking far more than it improves state tracking’s consequences. Anyone using auxiliary state supervision as a training signal, or probe accuracy as a progress metric, should treat that gap as the default expectation rather than the surprise.

7.10 A second domain, and what it took to build one

We attempted a second oracle-bearing domain, a program updating integer variables with interleaved queries, to test whether these findings are properties of long-horizon state tracking or of chess. Four designs failed to leave chance, and we report the sequence because the diagnosis is more useful than the intention.

The first placed a single query at the end of a sequence of random updates. Training loss settled at the token-type entropy of the generator, which is the value a model achieves by learning the format and nothing else, and the reason is arithmetic: the only learnable signal was one answer token in roughly one hundred and thirty, so state tracking carried under one percent of the gradient. Interleaving queries raised that share and did not move accuracy. Reducing the modulus and the number of variables did not move it either. A final version separated a pure read-out query from a computed one, so that state tracking could succeed without the arithmetic succeeding, and read-out accuracy still sat at chance.

We draw two conclusions from the arithmetic attempts. The narrow one is that this task family asks a small model to learn modular arithmetic before any state tracking becomes observable, and that our compute did not reach it. The broader one is a warning for anyone building synthetic state-tracking benchmarks: a task can be perfectly well specified, carry an exact oracle, and still measure nothing, because the quantity of interest contributes almost none of the training signal. The loss curve diagnoses this cheaply. A model that has learned the format and no content sits at the generator’s token-type entropy, which is computable in advance and worth computing before trusting such a benchmark. We report this in place of an easy cross-domain result.

A domain that does work. Abandoning arithmetic entirely, we built the second domain around an assignment instead. Four objects sit in four boxes. The sequence states the initial assignment, then applies operations that permute it, and after every operation asks where one object now is. A SWAP names two boxes and moves their contents without naming the objects, so the affected

object’s location cannot be retrieved by attending to its last mention and has to be maintained. A MOVE names the object, which gives an anchor and keeps the task learnable. The state is the assignment of objects to boxes, exact at every token, and the answer is a single symbol, so belief consistency here is exact rather than estimated: a wrong answer either is or is not the box the model’s own probe says holds that object.

Two further design choices were forced on us, and both are worth recording. Our first attempt used five objects. That makes the state an element of S_5 , a permutation-composition problem that is theoretically hard for the relevant fixed-depth model classes under asymptotic expressivity assumptions (Merrill et al., 2024). That may have contributed, but a theorem about asymptotic expressivity does not establish why one finite model failed to optimise on one bounded-length dataset, and we do not claim group structure as the cause. Four objects give S_4 , which is solvable. Our second attempt queried only occasionally, which left the task all-or-nothing: with uniformly random updates and sparse queries a model either tracks the permutation exactly or sits at chance, a structure that is hard to learn from endpoint supervision. Querying after every operation supplies partial credit, since a query one step after an operation is answerable directly, and that is the foothold from which the model extends. Chess has this property for free, because a plausible move needs only partial state.

With those two changes the task became measurable, and the outcome follows.

The outcome is not the clean replication we hoped for, and it is more informative than one would have been (Table 8).

With the supervised recurrent channel the task is close to solved at every depth, 0.989 at the deepest bucket against a chance rate of 0.25, the probe recovers the queried object at 0.983, and belief consistency is 0.701 against a mismatched control of 0.254. The mechanism reproduces.

The attention-only model behaves differently. It solves the task well above chance, 0.589 at depth, while the probe recovers its state at 0.263, which is chance. Belief consistency is 0.256 against a control of 0.249, which is no signal at all. The same probe procedure, on the same positions, recovers state at 0.983 in the other model, so this is a fact about the model and not a failure of the probe.

The reading we prefer is that this task does not force the model to maintain state. Only one object is asked about at a time, so the answer can be resolved on demand by tracing that object’s history, and an explicit assignment carried forward is unnecessary. Chess allows no such shortcut: predicting a legal move requires the whole position, so the state must be maintained eagerly.

That gives a fourth lesson for anyone building these diagnostics, and it is the one we would put first. A task can carry an exact oracle, be learnable, and still fail to measure state maintenance, because it never obliges the model to maintain anything. Belief consistency is only defined where a belief exists to read, and whether one exists is a property of the task, not only of the architecture.

7.11 Is the measurement useful, or only diagnostic?

Everything so far is a claim about mechanism. A reader is entitled to ask whether the same measurement is worth anything operationally, and the only baseline that matters is the model’s own confidence. If an external state monitor merely recovers what the model already knows about its own uncertainty, it is interesting and redundant.

We test this as selective prediction. A system that wants to avoid acting on a decayed state abstains on the positions it trusts least, and the question is what to rank by. We compare the model’s top-1 probability, the entropy of its next-move distribution, and the probe’s certainty about the squares the chosen move touches. That last signal uses only the probe’s own output, so it is computable at inference from activations, the probe having been fitted against an oracle

Table 12: Inference clustered on the game. Intervals resample whole games rather than positions, since positions within a game share an opening and a history. The lower block compares nested models by held-out log loss on a split by game, so each predictor’s contribution is measured out of sample. Lower is better.

Quantity	Estimate	95% CI (clustered)
belief consistency	0.3892	[0.3822, 0.3976]
mismatched-belief control	0.0654	[0.0620, 0.0687]
random-move control	0.0079	[0.0068, 0.0093]
AUC, whole board	0.6287	[0.6233, 0.6333]
AUC, move-touched	0.6961	[0.6926, 0.6998]
difference	+0.0674	[+0.0618, +0.0732]
Held-out log loss		
depth only	0.3630	
+ whole board	0.3595	
+ move-touched	0.3451	
+ both	0.3439	
+ both and confidence	0.3340	

offline.

One version of this experiment is not admissible and we report it only as a bound. Ranking by whether the probe is *wrong* on those squares requires the oracle at inference, which is the very thing an abstention rule exists to avoid. It is an upper bound on what a perfect state monitor could deliver, not a method.

The honest result is modest. Probe certainty alone is a weaker signal than the model’s own confidence, with an AUC of 0.593 against 0.628. Combined with confidence it is better than either alone: at half coverage the illegal-move rate among retained positions falls to 0.124 against 0.141 for confidence and 0.131 for entropy, and in a joint fit both terms carry weight, 0.554 for confidence and 0.367 for probe certainty. An external monitor therefore adds something the model does not self-report, and it adds less than we hoped.

The gap to the bound is the interesting part. Ranking by oracle-verified move-touched state reaches 0.093 at the same coverage, less than half the base rate of 0.197. Most of the benefit an ideal state monitor would deliver is not captured by the probe’s own certainty, which says the limitation is the quality of the state read-out rather than the premise. Better read-outs, not a different signal, are the lever.

7.12 Inference that respects the clustering

Positions within a game are not independent: they share an opening, a player, and an accumulating history. Intervals computed over positions therefore understate uncertainty, and every interval in this section resamples whole games instead. Section 7.5 already reports game-clustered intervals for belief consistency on the full evaluation population. This section is a restricted sensitivity analysis on a different population, 142,042 positions from 2,500 games of which 19,896 carry an illegal top-1 move, over ply 10 to 80 where the depth buckets used by the predictor comparison are defined. Estimates therefore differ slightly from the primary ones, and both are reported rather than reconciled by choosing whichever is larger.

Belief consistency is 0.389 [0.382, 0.398], against 0.065 [0.062, 0.069] for a mismatched belief and 0.008 [0.007, 0.009] for a random illegal move. The intervals are far apart, and the separation

does not depend on treating positions as independent.

The two predictors of an illegal move separate as well. Whole-board error reaches an AUC of 0.629 [0.623, 0.633] and move-touched error 0.696 [0.693, 0.700]. Bootstrapping the difference as a paired quantity gives +0.0674 [+0.0618, +0.0732], which excludes zero.

A population-averaged logistic fit, with game as the cluster and an exchangeable working correlation, gives coefficients with clustered standard errors on standardised predictors: +0.586 (clustered s.e. 0.011) for move-touched error against +0.161 (0.009) for the aggregate, with +0.573 for low model confidence and +0.633 for depth. The ordering survives clustered inference.

Out of sample, not only in sample. Coefficients can be significant and still add little. Splitting by game and fitting on seventy percent of games, the held-out log loss falls from 0.3630 with depth alone to 0.3595 adding whole-board error, and to 0.3451 adding move-touched error instead. Using both gives 0.3439, and adding the model’s own confidence 0.3340. The move-touched predictor buys several times what the aggregate buys out of sample, and it buys most of it even when the aggregate is already present.

A held-out abstention rule. Fitting and evaluating a selective prediction rule on the same positions overstates it. Refitting on training games and evaluating on held-out games, the illegal rate among retained positions at half coverage is 0.054 for a rule combining the model’s confidence with probe certainty, against 0.070 for confidence alone and 0.083 for probe certainty alone, from a base rate of 0.141. An oracle-verified state monitor would reach 0.038, which remains the bound rather than a method: establishing it at inference needs the ground truth the rule exists to avoid.

7.13 A constructive edit: moving a piece in the model’s belief

Suppressing a square’s occupancy removes the mass on moves from it, which is consistent with a causal belief state and equally consistent with generic damage to source-square features. A constructive edit separates the two. We take a square A the mover genuinely occupies and an empty square B , and edit toward “ A is empty” and “ B holds the piece that was on A ” together. Generic damage predicts that mass leaves A ; only a belief the policy reads predicts that mass arrives at B .

Both happen. Across 600 positions at the calibrated strength, mass on moves out of A changes by -0.0705 [-0.0810, -0.0603] while a norm-matched random edit gives -0.0052 [-0.0092, -0.0006]. Mass on moves out of B , which are illegal in the true position, changes by +0.0172 [+0.0134, +0.0222] against +0.0010 [+0.0006, +0.0016] for the random edit. Intervals are bootstrapped over games.

Splitting the edit budget across two directions makes the manipulation less reliable than the single-square version: A decodes as empty in 0.512 of cases and B as the transferred piece in 0.647. Restricting to the 205 cases where both flips demonstrably succeeded, the constructive effect strengthens, to +0.0291 [+0.0173, +0.0397], while the random control on A becomes indistinguishable from zero. We read this as the policy consulting a joined-up occupancy representation rather than a single feature, while noting it establishes this for occupancy and not for the whole position.

7.14 Off distribution: uniformly random legal play

Every measurement so far uses held-out human games, which are in distribution. Walker and Lyons (2026) report that a uniformly-random-legal-play split stays discriminative where in-distribution

Table 13: Transcript length with the position held fixed. A reversible cycle of legal moves returns the position exactly while adding four plies of history, so every row scores the same states at a different transcript length. Intervals are bootstrapped over games.

Added plies	occupied-square acc.	95% CI	illegal rate
0	0.736	[0.722, 0.747]	0.065
4	0.715	[0.702, 0.727]	0.050
8	0.706	[0.692, 0.718]	0.113
12	0.696	[0.683, 0.707]	0.177

performance saturates, so it is a fair test of whether our findings lean on familiar opening theory and typical structure.

They do not. On 1,200 games of uniformly random legal play, occupied-square accuracy falls from 0.889 to 0.112 and the illegal rate rises from 0.013 to 0.350, both steeper than on human games as expected for an unfamiliar distribution. Belief consistency is 0.380 over 32,027 illegal moves, against a random-move control of 0.008. That is slightly higher than the in-distribution rate, not lower. Coherent error is not an artifact of the model being on familiar ground.

7.15 Transcript length against position difficulty

The decay curves are read throughout as loss caused by composing more updates. Depth is confounded with position difficulty: a position at ply 100 differs from one at ply 10 in material, phase, and legal-move count. The curves alone do not license the causal reading, and this test removes the confound.

We insert a reversible cycle, a knight out and back for each side, which returns the position exactly while adding four plies of transcript. State is then read on an identical position at several transcript lengths.

Fidelity falls monotonically, from 0.736 [0.722, 0.747] with no insertion to 0.696 [0.683, 0.707] after twelve added plies, on 400 positions whose state is unchanged. The intervals at the two ends do not overlap. Transcript length alone therefore causes measurable state loss.

The size matters as much as the sign. Twelve plies of pure length cost about four points of accuracy, where the natural depth curve falls far further over its range. Both factors contribute, and we do not attribute the whole curve to composition depth. One caveat we cannot remove: the inserted moves are legal but unlike training data, so part of the drop may be distribution shift rather than length as such. The random-play result of Section 7.14 bounds that concern without eliminating it.

7.16 A policy-conditioned state-fidelity horizon

Both horizons used so far are read off behaviour, and the read-out ablation showed why that is fragile: a model can lower its illegal-move rate by hedging toward the legal-move prior, without knowing the position any better. A horizon defined on the representation is less exposed, because it does not read the confidence of the output: a model that spreads its mass more evenly gains nothing from doing so. It is not policy-independent, and we do not claim it is. The squares it scores are those the selected move touches, and that move comes from the output distribution, so a policy change can alter which squares are scored even with the state representation unchanged. It is a policy-conditioned state-fidelity horizon. A genuinely policy-independent comparison would

score a fixed set of components, for instance the recorded continuation move, every oracle-legal move averaged uniformly, or a candidate set shared across models.

We define H_{fid} as the depth at which the probability that the probe is wrong on the squares the model’s chosen move touches first crosses one half. It uses the same probe and the same activations as everything else, and it is immune to hedging: a model that spreads mass over legal moves gains nothing, since the measure asks what it knows, not what it emits.

Read this way the ladder is orderly. The smallest model reaches 33.8 ply, the baseline 40.8, and the largest 59.1. The behavioural horizon orders the same models the same way, 21.4, 26.1 and 33.8 ply, on a shorter scale because it asks a harder question of the model. Move-touched error at the shallowest depth runs from 0.207 in the smallest model to 0.078 in the largest. We report both horizons throughout, and prefer this one where the comparison is between architectures rather than between depths, since that is where hedging is most likely to confound a behavioural reading.

7.17 A degenerate metric, declared

Our horizon was pre-registered on exact-position fidelity, the probability that all sixty-four squares are simultaneously recovered. That criterion is already below its threshold in the very first bucket, so the pre-registered horizon evaluates to zero ply for every condition and cannot discriminate between them. The underlying quantity falls to 0.379 in the opening and to 0.000 by ply 20 to 30, after which exact states essentially do not occur, which is also why conditioning on them yields no samples at depth. We therefore report horizons on occupied-square accuracy and on the illegal-move rate. The revision was made after inspecting baseline data and before evaluating any intervention, and every definition is applied unchanged across conditions. We flag it rather than presenting the revised metric as the original plan, and we note that anyone measuring state tracking with an all-or-nothing criterion will obtain an unusable number for the same reason.

7.18 What the predictions came to

Pre-registration is only worth the words if the verdicts are stated, including the ones that went against us.

H1, supported. State fidelity declines with depth, from 0.975 to 0.066 occupied-square accuracy, clear of a majority predictor, a label-permutation control, and a randomised-weight model (Table 1, Figure 2). The behavioural horizon, 26.1 ply, sits well inside the context window.

H2, supported. The horizon grows slowly with scale across six rungs at three seeds each, from 21.4 ply at 3,440,064 parameters to 33.8 at 39,884,288 (Section 7.8).

H3, supported only weakly. The supervised recurrent channel raises state decodability substantially and the horizon by 26.5 against 26.1 ply. The architecture without supervision changes neither.

H4, not supported as predicted. We expected an intervention that improved state tracking to improve behaviour through it. Instead the intervention improved the measurement far more than its consequences, matching a model twice its size on decodability while capturing a fraction of the behavioural gain (Section 7.9). We designed H4 to be able to fail and it failed, in a direction that supports the paper’s main claim rather than undermining it.

H5, partially addressed. The natural-language transfer test our protocol specifies was not run and is absent rather than approximated (Section 9). A second oracle-bearing domain, reached only after several failed designs, tests the weaker claim (Section 7.10).

Three results were not predicted at all. The layer profile ran opposite to our expectation (Section 5). Ablating the recurrent read-out improved behaviour while worsening prediction, which

turned out to say something about our own observable rather than about the artifact (Section 7.9). And the standard whole-board measure of a world model proved a weaker predictor of behaviour than the same measure restricted to the state an action consumes, by a margin that grows with depth (Section 7.2).

If we could keep only one result it would be belief consistency together with the causal edit, which say that these models act coherently on a state they are losing. The measurement-placement result is a practical corollary: it tells you where to look, and it matters most exactly where long-horizon behaviour is hardest.

8 Discussion

The practical content of this paper is a change in what to measure. Reporting that a probe recovers a world model at some aggregate accuracy is weak evidence that the model is load-bearing: aggregate accuracy predicts the validity of the next action, and it predicts it poorly, and it gets worse with depth. Measuring the same representation on the components an action depends on is substantially more informative from the same activations. Any team instrumenting a system for long-horizon reliability should condition on the state the action uses.

The belief-consistency result reframes the failure. A model that plays a coherent move in a misremembered position is not behaving erratically, and treatments that assume erratic behaviour are aimed at the wrong target. Sampling several chains and voting may offer limited benefit when independently sampled continuations share the same state distortion, though we did not test repeated sampling or how far drift is correlated across samples. A verifier that scores outputs is well placed to catch the error and badly placed to prevent it, and it will pay that cost on every query for as long as the system runs.

The protocol generalises past chess. Anywhere an external checker can reconstruct the intended intermediate state, program execution traces, spreadsheet recalculation, structured database updates, ledger reconciliation, the same measurements are available. The requirement is not a game. It is an oracle.

We also treat the negative results as part of the contribution. A predicted layer profile came out backwards. A pre-registered metric turned out to be degenerate. A first design of a second domain turned out to be unlearnable, for a reason that was diagnosable from the loss curve. Each changed a decision, and each is reported where it did so.

What this implies for practice. The decomposition tells a team which remedy fits the failure they have. Where state-tracking error dominates, an inference-time verifier is an expensive way to buy back something the architecture gave away, and it is a recurring per-query cost rather than a one-off. Where planning error dominates, a verifier is exactly right. Reporting only the sum leaves a practitioner unable to tell which world they are in, and the two worlds recommend opposite purchases. The measurement needed to tell them apart is a rules engine and a linear probe, not privileged access to a model at scale, so it can be run in-house on a system a team already depends on. For the reader of a long generated answer, who is rarely in a position to audit step twenty-eight, the practical value is a stated depth beyond which the output should not be trusted without checking.

9 Limitations

- L1 Scale.** Our models are small and trained on a modest corpus, far below the chess-specific literature and any frontier system. Undertrained models are easier to study because they fail within a measurable range, which is what makes the instrument work, but the horizon values describe these models. Section 7.8 estimates the trend rather than assuming it away.
- L2 One domain family.** Chess carries an exact oracle by construction, which is what makes it measurable and also what makes it unrepresentative of open-ended reasoning. The second domain we built does reach a measurable regime (Section 7.10), but only its supervised condition maintains a decodable state, so the attention-only replication we wanted is not available there. The central findings rest on one instrument and should be read that way.
- L3 Probe decoding is lossy.** The believed board is an approximation, which is why we report the legal-move recovery rate as an explicit recovery reference rather than treating the decoded state as ground truth.
- L4 Pseudo-legality.** Belief consistency is tested with pseudo-legal moves, since a decoded board may lack a king and make check tests ill-defined. This is a permissive criterion, which the random-move control is designed to bound.
- L5 Scope of the causal claim.** The intervention establishes that the decoded state direction is read by the policy at the layer tested. It does not establish which downstream computation consumes it, and a single-direction edit is a coarse instrument.
- L6 A metric was revised after seeing data.** Declared in Section 7, applied uniformly, reported alongside the original.
- L7 No natural-language transfer test.** Establishing that these findings hold for mathematical or textual reasoning needs matched general-corpus pre-training, which our compute did not cover. It is absent rather than approximated.

10 Avenues of Future Research

- FR1** Extend the scale ladder far enough to estimate how horizon grows with parameters and data, and test whether the belief-consistency rate is scale-invariant even where the horizon moves.
- FR2** Replicate the decomposition in domains that were not designed as benchmarks, such as program execution or ledger reconciliation, where an oracle exists but the data is not constructed.
- FR3** Run the natural-language transfer test with matched pre-training, the experiment that decides whether this line of work reaches practice.
- FR4** Move from observation to intervention, editing move-touched state directions and measuring whether the induced action follows the induced belief.
- FR5** Develop state supervision that does not require an oracle, for example from execution traces or agreement between sampled continuations.
- FR6** Use move-touched fidelity as a live confidence signal, abstaining or escalating when the state an action depends on is poorly determined.
- FR7** Pre-register horizon metrics on pilot data, so metric selection is settled before confirmatory runs rather than declared afterwards.

11 Conclusion

Long-horizon behavioural failure can arise from degradation of the task state, from poor action selection given the state that is available, or from failures that remain unresolved under a lossy

decoder. These are not cleanly separable with the instruments we have, and we have not separated them.

What we can report is narrower and, we think, still worth having. In search-free chess models, whole-board probe fidelity is a weaker predictor of an invalid next action than fidelity on the endpoints of the move actually selected, and the gap widens with depth. A substantial identified subset of invalid actions is admissible under the board that model appears to hold at that position, at many times the rate of a mismatched-belief control and of an arbitrary invalid move. Targeted suppression and transfer edits show that the corresponding occupancy representation influences the policy rather than merely accompanying it. The identified share grows across the range of model sizes we can train, and survives a shift to uniformly random legal play.

These measurements identify state-consistent failure. They do not partition all failures, and the errors they leave unidentified are unresolved rather than attributed to reasoning. The practical consequence is a change in what to measure: a world model reported as an aggregate accuracy tells a practitioner less than the same representation scored on the state the next action will consume, and the difference matters most exactly where long-horizon behaviour is hardest to check.

Declarations

Disclosure statement. The authors report no conflict of interest.

Funding. This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Data availability statement. The training and evaluation code, the measurement pipeline, and the result files from which every reported value is generated are available at <https://github.com/MohitPratapSinghRathore/content-rot>.

AI usage declaration. During the preparation of this manuscript the authors used a generative AI tool to assist with language editing and refinement of the text. The authors reviewed and edited all output and take full responsibility for the content of the paper. Generative AI was not used to generate or analyse the study’s data, which were produced entirely by the evaluation pipeline described in the paper.

A Experimental detail

Data. Public Lichess archives, standard chess only, both players rated at least 1600, games between 30 and 160 ply, deduplicated by a hash of the move sequence. Moves are tokenised one per move in UCI form, giving a vocabulary of 1,926 including padding and a start symbol; there is no sub-move tokenisation and no board rendering in the input. Sequences are padded to 161 tokens and truncated at 160 ply. Splits are by game: 96,000 for language-model training, 12,000 for probe fitting, 12,000 held out for evaluation.

Oracle. Every game is replayed with a rules engine, recording after each token the occupancy of all 64 squares as one of 13 classes, the side to move, and castling rights. Replay completeness is verified before any model is measured.

Models. Pre-norm decoder-only Transformers with learned positional embeddings, GELU feed-forward of width four times the model dimension, and no dropout. The grid spans 6 to 12 layers and widths 192 to 512, with 6 or 8 heads. Training uses AdamW, $\beta = (0.9, 0.95)$, weight decay 0.01, a one-cycle schedule peaking at 3×10^{-4} with 5 percent warmup, batch size 32, 12000 steps, and mixed precision with gradient clipping at 1.0. The final checkpoint is used; no checkpoint selection is performed. Seeds 0, 1 and 2 control initialisation, batch order, and any auxiliary sampling.

Probes. One linear head per layer maps the residual stream to 64 squares by 13 classes, trained with cross-entropy on the probe split only, AdamW at 10^{-3} , two epochs, batch size 48. The reported layer is the one with the highest occupied-square accuracy, selected on the probe split rather than on the evaluation split. Controls are a per-square per-depth majority predictor, a label-permutation control pairing activations with another game’s state at equal depth, and a randomised-weight model of identical architecture whose probes are fitted the same way.

Believed boards. A decoded board is built by placing the argmax class at each square, setting side to move from ply parity, and clearing castling rights so that castling cannot be licensed by an assumption we did not decode. Legality is tested as pseudo-legality, since a decoded board may lack a king and make check tests ill-defined. Moves that fail to parse are counted as illegal.

Engine scoring. Stockfish 17 at fixed depth 10. Loss is the drop from the best available score to the score after the played move, both from the mover’s point of view, with mate scores mapped to 10000. The blunder threshold is 100 centipawns, fixed in advance.

Interventions. Edits are added to the residual stream at the output of the probed block, at the final position only. The direction is the difference of probe decoder rows for the target and current classes, normalised and scaled to α times the residual stream’s own root-mean-square magnitude times the square root of the model dimension. The reported α is 0.50, chosen from the sweep in Table 11 as the smallest value that reliably flips the decoded belief while leaving a norm-matched random edit at chance. Random controls are drawn isotropically and rescaled to the same norm.

Recurrent channel. A gated diagonal recurrence of width 64 inserted after every fourth block, with the read-out matrix zero-initialised so the module begins as the identity. The auxiliary objective, when active, is cross-entropy from the channel to the oracle state with weight 1.0. The parameter-matched control widens the feed-forward multiplier to 4.156 to absorb the same parameter count.

Second domain. Four objects in four boxes, an explicit initial assignment, then 24 operations drawn evenly between a swap of two named boxes and a move of a named object to a named box, with a query after every operation asking where one object is. Answers are weighted five times in the loss. Details of the six designs that failed before this one are in Section 7.10.

B Deviations from the pre-registration

We record every deviation rather than presenting the final protocol as the original one.

1. The horizon was pre-registered on exact-position fidelity, which proved degenerate. Horizons on occupied-square accuracy and on the illegal-move rate were added after inspecting baseline data and before evaluating any intervention, and applied unchanged to all conditions.

2. The bucketed horizon was replaced by an interpolated crossing point, because the bucketed read-off returned the same bucket for every condition.
3. Planning error was pre-registered as conditional on exact state recovery, which yields no samples past the opening. It is reported conditional on move-touched recovery, with a strictness ladder showing what the choice costs.
4. The aggregate and move-touched comparison was first made with incommensurable statistics, a correlation against a conditional probability. Both are now scored as predictors of the same outcome.
5. The layer-profile prediction was falsified and is retained as such.
6. The second domain required seven designs. All are reported.

C Research questions, design principles, and implications

This work began as a design-science exercise, and the framing is recorded here rather than in the body, where an empirical narrative reads better without it.

Research questions and objectives. **RQ1**, can long-horizon accuracy loss be separated into state tracking and planning, is addressed by **RO1**, defining both and giving an estimator for each (Sections 4 and 6). **RQ2**, does internal state loss explain behavioural failure or merely accompany it, is addressed by **RO2**, testing the link per token at the level of the state an action consumes (Sections 7.2 and 7.5). **RQ3**, how far does a search-free model reason before its state stops being recoverable and how does that respond to scale, is addressed by **RO3**, measuring horizons across a seeded ladder against trivial baselines (Section 7.8). **RQ4**, does the diagnostic discriminate between architectural interventions, is addressed by **RO4**, evaluating a recurrent state channel with parameter-matched and ablated controls (Section 7.9).

Design principles. **DP1**, ground the diagnostic in an external oracle rather than model self-report, realised by rules-engine labels at every token. **DP2**, separate the failure modes before attributing any effect, realised by scoring planning only where the move is legal and the relevant state is independently verified. **DP3**, establish what a trivial predictor achieves first, realised by majority, label-permutation and randomised-weight controls on every decodability number. **DP4**, measure the state the action depends on rather than the state in aggregate, realised by move-touched fidelity and its random-pair controls. **DP5**, test coherence by intervention and not only by correlation, realised by the calibrated suppression and transfer edits.

Implications for individuals, organisations, and society. For the reader of a long generated answer, who is rarely placed to audit step twenty-eight, the practical value is a stated depth beyond which output should not be trusted unchecked, and the uncomfortable finding that failure arrives fluent rather than obviously broken. For organisations, the decomposition indicates which remedy fits: a recurring per-query verifier cost suits a policy problem, a one-off structural change suits a memory problem, and reporting only the sum leaves a team unable to tell which they have. Strategically, the measurement needs a rules engine and a linear probe rather than privileged access at scale, so it can be owned and audited in-house rather than taken on trust. For society, diagnosing a failure is cheaper than paying inference-time compute to mask it, and remedies that cost a small fraction of parameters travel to on-device and public-sector deployments that cannot buy deliberation per request.

References

- Balogh, A. and Jelasi, M. (2026). Verification of the implicit world model in a generative model via adversarial sequences. In *International Conference on Learning Representations (ICLR)*.
- Gu, A. and Dao, T. (2024). Mamba: Linear-time sequence modeling with selective state spaces. In *First Conference on Language Modeling (COLM)*.
- Gu, A., Goel, K., and Ré, C. (2022). Efficiently modeling long sequences with structured state spaces. In *The Tenth International Conference on Learning Representations (ICLR)*.
- Harang, R., Naradowsky, J., Gujju, S. V., and Miyao, Y. (2025). Tracking world states with language models: State-based evaluation using chess. In *ICML Workshop on Assessing World Models*. Spotlight; arXiv:2508.19851.
- Hewitt, J. and Liang, P. (2019). Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2733–2743. Association for Computational Linguistics.
- Hsieh, C.-P., Sun, S., Krizan, S., Acharya, S., Rekesh, D., Jia, F., Zhang, Y., and Ginsburg, B. (2024). RULER: What’s the real context size of your long-context language models? In *First Conference on Language Modeling (COLM)*.
- Karvonen, A. (2024). Emergent world models and latent variable estimation in chess-playing language models. In *First Conference on Language Modeling (COLM)*.
- Karvonen, A., Wright, B., Rager, C., Angell, R., Brinkmann, J., Smith, L., Verdun, C. M., Bau, D., and Marks, S. (2024). Measuring progress in dictionary learning for language model interpretability with board game models. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*.
- Li, K., Hopkins, A. K., Bau, D., Viégas, F., Pfister, H., and Wattenberg, M. (2023). Emergent world representations: Exploring a sequence model trained on a synthetic task. In *The Eleventh International Conference on Learning Representations (ICLR)*.
- Li, Q. and Jiang, W. (2026). Tracking vs. deciding: The dual-capability bottleneck in searchless chess transformers. arXiv:2603.29761.
- Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., and Cobbe, K. (2024). Let’s verify step by step. In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. (2024). Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173.
- Meng, K., Bau, D., Andonian, A., and Belinkov, Y. (2022). Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*.
- Merrill, W., Petty, J., and Sabharwal, A. (2024). The illusion of state in state-space models. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235, pages 35492–35506.

- Merrill, W. and Sabharwal, A. (2024). The expressive power of transformers with chain of thought. In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Nanda, N., Lee, A., and Wattenberg, M. (2023). Emergent linear representations in world models of self-supervised sequence models. In *Proceedings of the 6th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 16–30. Association for Computational Linguistics.
- Pereira, D. and Zuidema, W. (2026). Transformers struggle to use their emergent world models: Revisiting the tower of hanoi, and the illusion of thinking. arXiv:2608.07077.
- Ruoss, A., Delétang, G., Medapati, S., Grau-Moya, J., Wenliang, L. K., Catt, E., Reid, J., Lewis, C. A., Veness, J., and Genewein, T. (2024). Amortized planning with large-scale transformers: A case study on chess. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*.
- Toshniwal, S., Wiseman, S., Livescu, K., and Gimpel, K. (2022). Chess as a testbed for language model state tracking. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 11385–11393.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems 30 (NIPS 2017)*.
- Walker, B. and Lyons, T. (2026). Chess-world-model: A 10m-game benchmark for exact state tracking from chess move sequences. arXiv:2605.30100.
- Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E. H., Narang, S., Chowdhery, A., and Zhou, D. (2023). Self-consistency improves chain of thought reasoning in language models. In *The Eleventh International Conference on Learning Representations (ICLR)*.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E. H., Le, Q. V., and Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*.
- Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T. L., Cao, Y., and Narasimhan, K. (2023). Tree of thoughts: Deliberate problem solving with large language models. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*.